

PrivacyScrub – Software Requirements Specification (V4)

1. Introduction

1.1 Purpose

This document defines the **V4 Software Requirements** for **PrivacyScrub**, an API-first media anonymization platform that detects and obscures sensitive information in images and videos to support privacy and regulatory compliance (GDPR, CCPA, HIPAA Safe Harbor–inspired) for user-generated and professional media.

V4:

- **Includes and refines all requirements from V1–V3**, including:
 - * Image & video anonymization.
 - * Asynchronous job handling, chunking, and stitching.
 - * Local inference (no per-frame external API calls).
 - * GPU acceleration and hardware video codecs.
 - * Compliance profiles and evaluation pipeline.
- Introduces a **modern, multi-model anonymization stack** and world-class MLOps/production guarantees.

This SRS is intended for:

- Backend engineers (FastAPI, async workflows, GCP).
- ML engineers (detection models, evaluation, active learning).
- DevOps/SRE/MLOps (deployment, observability, GPU infra).
- Security, privacy, and compliance stakeholders.

1.2 Scope

PrivacyScrub V4 provides:

- API-first anonymization for **images and videos**.
- Detection and redaction of:
 - * Faces
 - * License plates
 - * Logos / brand marks
 - * On-screen text (signs, jerseys, documents, name labels)
 - * EXIF and container metadata (e.g., GPS, device IDs, author fields)
 - * Other visual PII/PHI relevant to GDPR/CCPA/HIPAA principles
- Multiple anonymization modes:
 - * Blur
 - * Pixelate
 - * Black box (solid mask)
- Asynchronous video processing with chunking and parallel processing.
- Local, GPU-accelerated inference stack (no frame-level HTTP calls).

- Compliance profiles: * **NONE** * **GDPR** * **CCPA** * **HIPAA_SAFE_HARBOR**
- Evaluation pipeline with public benchmarks and internal test suites.
- Streamlit (or similar) UI as a reference client for the APIs.

Out of scope for V4, but considered future work:

- Audio anonymization (voice redaction, beeps, etc.).
- Full document anonymization beyond text within images/videos.
- On-premise and air-gapped deployments (architecture should not preclude them).

1.3 Definitions and Acronyms

- **PII** – Personally Identifiable Information
- **PHI** – Protected Health Information
- **GDPR** – General Data Protection Regulation (EU)
- **CCPA** – California Consumer Privacy Act
- **HIPAA Safe Harbor** – HIPAA de-identification method based on removing or generalizing 18 identifiers
- **ROI** – Region of Interest (rectangular region in image/video frames)
- **RT-DETR** – Real-Time Detection Transformer (or similar efficient general object detector)
- **DBNet / CRAFT / EAST** – State-of-the-art scene text detection architectures
- **Logos in the Wild** – Example logo detection dataset for logo anonymization

2. System Overview

2.1 Product Vision

PrivacyScrub allows users—from journalists to real estate agents to marketers—to:

“Upload media → configure compliance/profile → receive anonymized media or detection masks → publish confidently without exposing unnecessary PII/PHI or brand conflicts.”

V4’s vision is to be a **world-class anonymization engine**, with:

- High recall and precision for critical PII/PHI elements.
- Real-time or faster-than-real-time video throughput in GPU deployments.
- Compliance-aware defaults via profiles.
- Measurable performance and quality via standardized evaluation.

2.2 Architecture Perspective

Core components:

1. **API Service (FastAPI on Cloud Run/GKE)** - REST endpoints for image & video anonymization. - Job creation, job status, cancellation endpoints. - Stateless; uses durable backing services for state.
2. **Job Orchestrator** - Cloud Tasks (or equivalent) for chunk-based video processing. - Firestore (or equivalent) for job metadata: - Status - Chunk information - Options snapshot - Metrics and errors
3. **Worker Service (GPU-capable)** - Local inference engine with a **multi-model stack**: - RT-DETR (or similar) for general object detection (people, vehicles, etc.). - Dedicated face detection model (e.g., RetinaFace/YOLO-face). - Dedicated license plate detection model. - Scene text detection model (DBNet/CRAFT/EAST style) plus text recognizer. - Logo detection model fine-tuned on a logo dataset (e.g., Logos in the Wild). - Hardware-accelerated video decode/encode (NVDEC/NVENC or equivalent). - Anonymization (blur/pixelate/black box) applied frame-wise.
4. **Storage (Google Cloud Storage)** - Temporary storage of: - Original inputs - Chunked intermediates - Final anonymized outputs - Strict lifecycle policies and TTL-based deletion.
5. **Model Registry & Configuration** - Versioned records of model artifacts and config: - Model names and hashes - Configs per target type (faces, plates, logos, text) - Integration with evaluation pipeline and deployment gating.
6. **Frontend UI (Streamlit)** - Not authoritative for application logic. - Convenience client for demos, QA, and manual validation. - Uses only documented API endpoints.

3. Functional Requirements

All requirements from V1–V3 are retained and superseded where necessary. V4 adds a multi-model stack, improved compliance, and production hardening.

3.1 Targets & Compliance Profiles

3.1.1 Target Categories

Category 1 – Core PII (Always Supported)

- Faces (full and identifiable partial faces).
- License plates (all regional formats).
- EXIF/metadata: * GPS coordinates * Device IDs / camera model and serial * Timestamps (where appropriate) * Author / creator fields
- Household-level visual identifiers: * Nameplates on doors/mailboxes * Name badges (where detectable)

Category 2 – Toggleable High-Value Targets

- Logos / brand marks: * Clothing logos * Store signs * Product packaging

- On-screen text: * Street signs * Store names * Jersey names and numbers * Document text
- Financial information: * Credit/debit card numbers * Bank account numbers
- Signatures: * Handwritten or digital signatures
- Tattoos / strong unique identifying marks
- Other high-risk text: * Email addresses * Phone numbers * Physical addresses

Category 3 – Advanced / Location-Based

- Specific addresses and location details from visible signs.
- Highly distinctive building facades, when combined with strong location cues.
- Contextual identifiers: * Hospital or clinic names (medical context) * School names in footage featuring minors.

For V4, Category 3 is handled via a combination of text recognition, metadata removal, and object detection. Full “scene uniqueness” modeling remains a future enhancement.

3.1.2 Compliance Profiles

`compliance_profile: "NONE" | "GDPR" | "CCPA" | "HIPAA_SAFE_HARBOR"`

NONE

- Defaults: * `\text{faces} = \text{true}` * `\text{plates} = \text{true}` * `\text{logos} = \text{true}` * `\text{text} = \text{true}`
- `\text{mode} = \text{blur}`
- `\text{confidence_threshold} = 0.5`
- `\text{strip_metadata} = \text{true}`
- User may toggle any target on or off, and change mode/threshold.

GDPR

- Force-enable: * Faces * License plates * Text related to names and addresses
- `\text{strip_metadata} = \text{true}` (especially GPS and device IDs).
- Minimum `\text{confidence_threshold}` (e.g., `\geq 0.6`).
- Default `\text{mode} = \text{blur}`.
- User cannot disable core PII targets; may enable additional ones (logos, tattoos, etc.).

CCPA

- Force-enable: * Faces * License plates * Text relating to household-level identifiers (addresses, family names)
- `\text{strip_metadata} = \text{true}`.
- Similar or slightly higher thresholds than default.
- User cannot disable core PII targets; may enable additional ones.

HIPAA_SAFE_HARBOR

- Force-enable:
 - * Faces (full-face and comparable images)
 - * License plates / vehicle identifiers
 - * Text that contains:
 - * Names
 - * Contact info
 - * Account numbers
- `$\text{strip_metadata} = \text{true}` (GPS and other potential identifiers).
- Default `$\text{mode} = \text{black_box}` or strong `\text{pixelate}`.
- Minimum `$\text{confidence_threshold} \geq 0.7`.
- Documented as a technical aid only; no legal guarantee of compliance.

FR-CP-01: Both image and video paths MUST apply `$\text{get_config_for_profile}` and enforce these constraints.

FR-CP-02: User options MAY further restrict targets (turn off non-mandatory ones), but MUST NOT disable profile-mandated targets.

3.2 Image Anonymization (Synchronous)

Endpoint: `POST /v1/anonymize-image`

FR-IMG-01: Accepts:

- File: `$\text{image/jpeg}` or `$\text{image/png}` as `$\text{multipart/form-data}`.
- Options via form fields or nested JSON, including at least:
 - * `$\text{targets.faces: bool}`
 - * `$\text{targets.plates: bool}`
 - * `$\text{targets.logos: bool}`
 - * `$\text{targets.text: bool}`
 - * `$\text{mode: blur | pixelate | black_box}`
 - * `$\text{confidence_threshold: float}`
 - * `$\text{coordinates_only: bool}`
 - * `$\text{compliance_profile: string}` * Optional `$\text{roi: [x1, y1, x2, y2]}` in normalized or pixel coordinates.

FR-IMG-02: The endpoint MUST:

1. Parse the form fields / JSON.
2. Construct a `$\text{PrivacyConfig}` from the compliance profile.
3. Override `$\text{PrivacyConfig}` with user-specified options while respecting mandatory profile protections.
4. Run detection on the image using the V4 detection stack (see [S 3.4](#)).
5. Apply anonymization within detected bounding boxes, respecting:
 - Selected targets
 - Mode
 - Confidence threshold
 - ROI
6. Return:
 - An anonymized image (binary) if `$\text{coordinates_only} = \text{false}`, or
 - A JSON structure describing detections if `$\text{coordinates_only} = \text{true}`.

FR-IMG-03: If `$\text{strip_metadata} = \text{true}`, output image MUST have all `$\text{EXIF/metadata}` removed.

FR-IMG-04: The API MUST return structured JSON error payloads for invalid input (type, size, corrupted file, invalid options).

3.3 Video Anonymization (Asynchronous, Chunked)

Endpoints:

- \$text{POST /v1/anonymize-video} \rightarrow\$ submit a video for anonymization.
- \$text{GET /v1/jobs/{job_id}} \rightarrow\$ get job status and info.
- \$text{DELETE /v1/jobs/{job_id}} \rightarrow\$ cancel a job.

3.3.1 Job Lifecycle

FR-VID-01: Job statuses:

- \$text{QUEUED}\$ – Job metadata created, video uploaded.
- \$text{CHUNKING}\$ – Splitting video into chunks.
- \$text{PROCESSING}\$ – Chunks being processed.
- \$text{STITCHING}\$ – Processed chunks being combined.
- \$text{COMPLETED}\$ – Final video available.
- \$text{FAILED}\$ – Error occurred; error message should be provided.
- \$text{CANCELLED}\$ – Job cancelled by client or system.

FR-VID-02: \$text{POST /v1/anonymize-video}\$ MUST:

- Validate the uploaded video (e.g., \$text{video/mp4}\$ \$text{v1}\$).
- Upload it to GCS (e.g., \$text{input/{job_id}/original.mp4}\$).
- Create a Firestore job document with: * \$text{status = "QUEUED"}\$ * \$text{chunks_total} = 0\$ * \$text{chunks_completed} = 0\$ * Options snapshot (including \$text{profile}\$, \$text{targets}\$, \$text{mode}\$, \$text{threshold}\$).
- Enqueue a split task (e.g., \$text{/internal/split-video}\$).

3.3.2 Splitting & Chunk Processing

FR-VID-03: Split task (\$text{/internal/split-video}\$) MUST:

- Inspect video duration.
- If \$text{duration} \le \text{MIN_CHUNK_DURATION_SEC}\$: * Set \$text{chunks_total} = 1\$. * Use original file as the single chunk.
- Else: * Use \$text{ffmpeg}\$ to segment into \$text{N}\$ chunks (default \$sim 5\$ minutes). * Store each chunk at \$text{input/{job_id}/chunks/chunk_{i}.mp4}\$. * Set \$text{chunks_total} = N\$.
- Set job \$text{status} = "CHUNKING\$".
- Enqueue one process-chunk task per chunk.

FR-VID-04: Process-chunk task (\$text{/internal/process-chunk}\$) MUST:

- For the first started chunk, update job status from \$text{CHUNKING} \rightarrow \$text{PROCESSING} (if not already updated).

- Download the chunk file from GCS to local disk.
- Run anonymization frame-wise using the V4 detection engine (see \$S 3.4\$).
- Use a writeable frame copy to avoid “assignment destination is read-only” errors.
- Encode the processed chunk to `$text{output/{job_id}/chunks/chunk_{i}.mp4$}` with hardware acceleration if available.
- Upload processed chunk to GCS.
- Increment `$text{chunks_completed$}` in Firestore.
- If `$text{chunks_completed} \geq text{chunks_total$}` and no chunk has failed: * Enqueue `$text{/internal/stitch-video$}`.

3.3.3 Stitching & Completion

FR-VID-05: Stitch task (`$text{/internal/stitch-video$}`) MUST:

- Set job status to `$text{STITCHING$}`.
- Concatenate processed chunks in correct order using `$text{ffmpeg$}`.
- Strip metadata from final video output.
- Write final video to `$text{output/{job_id}/final.mp4$}`.
- Generate a signed URL or API-accessible download path.
- Update job: * `$text{status = "COMPLETED"$}` * `$text{output_url$} = <text{generated_url}>$`

3.3.4 Error Handling & Cancellation

FR-VID-06: Any unhandled exception in split/process/stitch tasks MUST:

- Set job `$text{status = "FAILED$}`.
- Populate `$text{error_message$}` with a concise description.

FR-VID-07: `$text{DELETE /v1/jobs/{job_id}$}` MUST:

- Set job `$text{status = "CANCELLED$}`.
- Prevent scheduling of further tasks for that job.
- Workers MUST check for cancellation before heavy work and abort gracefully when safe.

3.3.5 Progress Reporting

FR-VID-08: `$text{GET /v1/jobs/{job_id$}}` MUST return:

- `$text{job_id$}`
- `$text{status$}`
- `$text{created_at$}, $text{updated_at$}`
- `$text{chunks_total$}`
- `$text{chunks_completed$}`
- Optional `$text{chunks_failed$}`

- Derived \$text{progress} (\$text{0.0 – 1.0}) when relevant
- \$text{output_url} and \$text{TTL} (if \$text{COMPLETED})
- \$text{error_message} (if \$text{FAILED})

3.4 Detection & Anonymization Engine (V4 Model Stack)

V4 introduces a multi-model detection stack instead of a single \$text{YOLO-style} model.

3.4.1 Detection Interface

FR-DET-01: All detection MUST flow through a central interface:

```
def detect_frame(
    frame: np.ndarray,
    config: PrivacyConfig
) -> List[Detection]:
    """
```

Returns a list of detections, each with:

- type (face, plate, logo, text, tattoo, etc.)
- bbox/poly coordinates
- confidence
- metadata (e.g., recognized text)

"""

FR-DET-02: The detector MUST orchestrate the following internal sub-models:

- **General Object Detector:** Model: \$text{RT-DETR} or equivalent real-time detector.
- **Face Detector:** Dedicated face detection model (e.g., \$text{RetinaFace/YOLO-face}).
- **License Plate Detector:** Dedicated license plate detection model.
- **Scene Text Detection + OCR:** Text detection model (\$text{DBNet/CRAFT/EAST} style) plus text recognizer.
- **Logo Detection:** Logo-specific detector fine-tuned on a logo dataset.

FR-DET-03: Detection outputs MUST include:

- \$text{type} (\$text{e.g., "face", "plate", "logo", "text", "tattoo"})
- \$text{bbox} and/or \$text{poly} (\$text{coordinates})
- \$text{confidence}
- Optional: \$text{text content} (\$text{for text detections})

FR-DET-07: Implementation MUST operate on a writeable frame array:

- At start of frame pipeline: \$text{frame = np.array(frame, copy=True)} or equivalent.
- Avoid in-place modifications to read-only views.

3.5 Coordinates-Only Mode & ROI

FR-COORD-01: If `$text{coordinates_only} = true$`:

- **Images** (`$text{POST}$`): Returns a `$text{JSON payload directly$}` with detections.
- **Videos** (`$text{GET}$`): The worker writes a `$text{JSON manifest$}` to `$text{GCS$}` with per-frame detections; `$text{GET /v1/jobs/{job_id}}$` returns the `$text{GCS$}` link to this manifest in `$text{output_url$}`.

ROI Handling (`$text{FR-IMG-01$} / $text{V4 Detection$}`):

- The detection functions accept an optional `$text{ROI$}` parameter.
- For **images** and **video frames**, only detections whose bounding boxes intersect the `$text{ROI$}` rectangle are considered for anonymization.

3.6 Reference UI (Streamlit)

The Streamlit UI acts as a thin client over the public API (`$text{FR-UI-01$}`).

- **Images:** Provides upload controls and `$text{UI$}` elements for selecting targets, `$text{mode$}`, `$text{profile$}`, `$text{confidence$}`, and optional `$text{ROI$}` inputs.
- **Videos:** Calls `$text{POST /v1/anonymize-video$}` for submission, then polls `$text{GET /v1/jobs/{job_id$}}` to display status, chunk progress, and the final `$text{output_url$}`.
- **QA View (FR-HIL-01):** An internal `$text{QA$}` mode renders original vs. anonymized frames side-by-side with detection overlays for review. Reviewer adjustments are stored in an access-controlled dataset for future model retraining and analysis.

4. Non-Functional Requirements (NFR)

4.1 Performance and GPU Acceleration (NFR-PERF-01)

- **Hardware Codecs:** The worker utilizes `$text{NVDEC/NVENC$}` (via `$text{ffmpeg$}`) for hardware video decode and encode to achieve the `$ge 1{text{x$}` real-time `$text{SLO$}`.
- **Resource Allocation:** `$text{Cloud Run$}` is configured with `$text{4 GiB$}` memory and `$text{2 vCPUs$}` to support the multi-model stack.

4.2 Security and Observability

- **Logging Hygiene (NFR-SEC-03):** Logs MUST be scrubbed to avoid storing raw frames or PII. Only job IDs, technical metrics, and redacted error messages are permitted.
- **Metrics and CI/CD Gating (FR-EVAL-01 / FR-EVAL-03):**
 - An automated evaluation suite runs against benchmark datasets (`$text{faces, plates, logos, text$}`) to compute metrics (`$text{F1, mAP$}`).

- \$text{CI/CD}\$ promotion is gated; new models MUST pass configured metric thresholds.
- **Drift Monitoring:** Detection statistics (\$text{confidences, targets-per-frame}\$) are periodically tracked and compared against baselines.
- **Data Retention (NFR-SEC-04):** GCS bucket lifecycle rules MUST enforce \$text{TTL}\$ deletion of all media files.
- **Multi-Tenancy:** The \$text{API}\$ gateway enforces per-tenant quotas and rate limits, and metrics are segmented by tenant identifier.