

PrivacyScrub V4: Technical Design Document (TDD)

Author: Greg Burns **Version:** 4.2 (Final Design Document) **Date:** November 20, 2025

1. Introduction and Project Goals

This TDD outlines the technical blueprint for PrivacyScrub V4. The document is designed to achieve **100% compliance** with the V4 SRS, leveraging a local, GPU-accelerated, multi-model architecture for video processing.

1.1 V4 Implementation Mandate

The design adheres to these core architectural principles:

- **Local Model Fusion:** Eliminate network latency by hosting the entire multi-model stack (Faces, Plates, Text, Logos) locally on the worker.
- **Compliance Enforcement:** Rigorously enforces non-negotiable profile constraints (\$\text{GDPR}\$, \$\text{HIPAA}\$) within the configuration logic (\$\text{get_config_for_profile}\$).
- **Stability:** Implements explicit fixes for the known V3 pipeline bugs (\$\text{read-only array}\$ and \$\text{stuck status}\$).

2. System Architecture

The V4 architecture is a high-performance, distributed, event-driven model on GCP, optimized for high concurrency.

2.1 Worker Service: V4 Multi-Model Detection Stack

The worker container hosts the entire model stack and the processing pipeline, fulfilling \$\text{FR-DET-02}\$.

- **Backbone / General:** \$\text{RT-DETR}\$ or \$\text{YOLOv8}\$ (Local). Provides general object context (Persons, Vehicles).
- **Face Detector:** Dedicated Face model (Local). Provides high-accuracy detection of faces.
- **Plate Detector:** Dedicated Plate model (Local). Provides high-accuracy license plate identification.

- **Text Detector:** Local \$\text{Scene Text Detector}\$ (\$\text{DBNet/CRAFT}\$) plus \$\text{OCR Recognizer}\$. Extracts bounding boxes and content for sensitive text.
- **Logo Detector:** Local \$\text{Logo Classifier}\$ (e.g., fine-tuned YOLO head). Identifies high-value brand marks.
- **Fusion Layer:** Custom Python logic combining bounding box results, resolving conflicts, and applying \$\text{ROI}\$ filtering before anonymization.

2.2 Job Lifecycle and Endpoints

The API surface defines the contract for consuming the asynchronous worker service.

- **Endpoints:**
 - \$\text{POST /v1/anonymize-video}\$: Initiates job.
 - \$\text{GET /v1/jobs/\{job_id\}}\$: Returns the full status object.
 - \$\text{DELETE /v1/jobs/\{job_id\}}\$: Sets \$\text{status}=\text{CANCELLED}\$ (FR-VID-07).
- **Status Transitions (FR-VID-01):** The worker implements explicit state changes to guarantee accurate client feedback:
 - \$\text{QUEUED} \rightarrow \text{CHUNKING} \rightarrow \text{PROCESSING} \rightarrow \text{STITCHING} \rightarrow \text{COMPLETED}\$.

3. Functional Requirements and Implementation Details

3.1 Job Lifecycle API Endpoints (FR-VID-01, FR-VID-07, FR-VID-08)

The following public endpoints expose the asynchronous video processing lifecycle:

- \$\text{POST /v1/anonymize-video}\$
 - Validates the uploaded video (type/size).
 - Uploads to GCS under \$\text{input}/\{job_id\}/original.mp4\$.
 - Creates a Firestore job document with \$\text{status = "QUEUED"}\$ and captured configuration.
 - Enqueues the \$\text{/internal/split-video}\$ task.
- \$\text{GET /v1/jobs/\{job_id\}}\$
 - Returns a status object containing:
 - \$\text{status}\$ (\$\text{QUEUED}\$, \$\text{CHUNKING}\$, \$\text{PROCESSING}\$, \$\text{STITCHING}\$, \$\text{COMPLETED}\$, \$\text{FAILED}\$, \$\text{CANCELLED}\$)
 - \$\text{chunks_total}\$, \$\text{chunks_completed}\$ (for progress)
 - \$\text{output_url}\$ (if \$\text{COMPLETED}\$ or coordinates-only manifest ready)
 - \$\text{error_message}\$ (if \$\text{FAILED}\$).
- \$\text{DELETE /v1/jobs/\{job_id\}}\$
 - Sets \$\text{status = "CANCELLED"}\$ for the job.

- Workers MUST check for `$\text{CANCELLED}` before starting heavy work and abort gracefully if set.

3.2 Compliance Profile Enforcement (FR-CP-02)

The configuration logic strictly enforces profile requirements, overriding user inputs where necessary.

- **GDPR:** `$\text{confidence_threshold} \geq 0.6$, $\text{target_faces} = \text{True}$, $\text{target_text} = \text{True}$, $\text{strip_metadata} = \text{True}$.`
- **HIPAA_SAFE_HARBOR:** `$\text{confidence_threshold} \geq 0.7$, $\text{mode} = \text{BLACK_BOX}$, $\text{strip_metadata} = \text{True}$, $\text{target_all} = \text{True}$.`
- **CCPA:** `$\text{confidence_threshold} \geq 0.55$, $\text{target_faces} = \text{True}$, $\text{target_plates} = \text{True}$, $\text{strip_metadata} = \text{True}$.`
- **NONE:** Default profile. `$\text{target_all} = \text{True}$, $\text{mode} = \text{BLUR}$, $\text{confidence_threshold} = 0.5$, $\text{strip_metadata} = \text{True}$.`

3.3 Detection Interface and Data Output

- **Image API Wiring (FR-IMG-01):** The `$/v1/anonymize-image$` endpoint accepts all `Form` fields for per-target toggles (`$\text{target_faces}`, `$\text{target_plates}`, `$\text{target_logos}`, `$\text{target_text}`) and `$\text{mode}`, applying profile constraints first.
- **EXIF / Metadata Strip (FR-IMG-03):** For images, when `$\text{strip_metadata} = \text{True}$` (e.g., under `$\text{GDPR}`), the `$/v1/anonymize-image$` implementation MUST re-encode the output using a clean `$\text{Pillow}` save path, ensuring all `$\text{EXIF}` and other metadata are removed from the returned file.
- **Coordinates-Only Mode (FR-COORD-01):**
 - **Images** (`$\text{POST}`): Returns a `$\text{JSON payload directly}` with detections.
 - **Videos** (`$\text{GET}`): The worker writes a `$\text{JSON manifest}` to `$\text{GCS}` with per-frame detections; `$/v1/jobs/{job_id}$` returns the `$\text{GCS}` link to this manifest in `$\text{output_url}`.
- **ROI Handling (FR-IMG-01):** The `$\text{Model Fusion Layer}` filters all detections to ensure only bounding boxes that intersect the optional `$\text{ROI}` rectangle are considered for anonymization in both images and video frames.

3.4 Reference UI (Streamlit)

The Streamlit UI acts as a thin client over the public API.

- **Images:** Provides configuration controls (targets, mode, profile, `$\text{ROI}`) and renders a before/after preview.

- **Videos:** Calls `POST /v1/anonymize-video` for submission, then polls `GET /v1/jobs/{job_id}` to display status, chunk progress, and the final `output_url`.
- **QA View:** An internal `QA` mode renders original vs. anonymized frames side-by-side with detection overlays for review (`FR-HIL-01`).

4. Operational and MLOps Requirements

4.1 Performance and GPU Acceleration (NFR-PERF-01)

- **Frame Mutability (FR-DET-07):** The `frame_processor` MUST create a writeable copy immediately (`frame = np.array(frame, copy=True)`) to eliminate the V3 "read-only array" bug.
- **Hardware Codecs:** The worker utilizes `NVDEC/NVENC` (via `ffmpeg`) for hardware video decode and encode to achieve the `ge 1x` real-time SLO.
- **Resource Allocation:** `Cloud Run` is configured with `4 GiB` memory and `2 vCPUs` to support the multi-model stack.

4.2 Security and Observability

- **Logging Hygiene (NFR-SEC-03):** Logs MUST be scrubbed to avoid storing raw frames or PII. Only job IDs, technical metrics, and redacted error messages are permitted.
- **Metrics and CI/CD Gating (FR-EVAL-01 / FR-EVAL-03):**
 - An automated evaluation suite runs against benchmark datasets (`faces, plates, logos, text`) to compute metrics (`F1, mAP`).
 - `CI/CD` promotion is gated; new models MUST pass configured metric thresholds.
- **Drift Monitoring:** Detection statistics (`confidences, targets-per-frame`) are periodically tracked and compared against baselines to flag potential model degradation.
- **Data Retention (NFR-SEC-04):** GCS bucket lifecycle rules MUST enforce `TTL` deletion of all media files.
- **Multi-Tenancy:** The `API` gateway enforces per-tenant quotas and rate limits, and metrics are segmented by tenant identifier.